

In [7]:

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from collections import Counter
import lightgbm as lgb
import seaborn as sns
```

In [8]:

```
train = pd.read_csv("/Users/BarryFitzpatrick/Machine Learning/Kaggle Group/tcd-ml-comp-201920-income-pred-group/tcd-ml-1920-group-income-train.csv")
test = pd.read_csv("/Users/BarryFitzpatrick/Machine Learning/Kaggle Group/tcd-ml-comp-201920-income-pred-group/tcd-ml-1920-group-income-test.csv")
```

```
train = train.sample(frac = 1)
train.shape
```

```
/Users/BarryFitzpatrick/anaconda3/lib/python3.7/site-packages/IPython/core/interactiveshell.py:2785: DtypeWarning: Columns (2,4) have mixed types. Specify dtype option on import or set low_memory=False.
  interactivity=interactivity, compiler=compiler, result=result)
/Users/BarryFitzpatrick/anaconda3/lib/python3.7/site-packages/IPython/core/interactiveshell.py:2785: DtypeWarning: Columns (4) have mixed types. Specify dtype option on import or set low_memory=False.
  interactivity=interactivity, compiler=compiler, result=result)
```

Out[8]:

(1048574, 17)

In [9]:

```
train_missing = (train.isnull().sum()/len(train))*100
train_missing = train_missing.drop(train_missing[train_missing==0].index).sort_values(ascending=False)
miss_data = pd.DataFrame({'缺失百分比':train_missing})
miss_data
```

Out[9]:

	缺失百分比
University Degree	7.686630
Gender	7.069315
Hair Color	6.695856
Satisfaction with employer	3.632266
Year of Record	0.382710
Profession	0.272084

In [5]:

```
train.head()
```

Out[5]:

Instance	Year of Record	Housing Situation	Crime Level in the City of Employment	Work Experience in Current Job [years]	Satisfaction with employer	Gender	Age	Country	Size of City	Profession	Unive De
----------	----------------	-------------------	---------------------------------------	--	----------------------------	--------	-----	---------	--------------	------------	----------

701573	644709	1992.0	Small House	90	10	Average	unknown	17	Dominican Republic	1627279	project manager	
167938	Instance 167939	Year of Record 1953.0	Housing Situation nA	Crime Level in the City of Employment 23	Work Experience in Current Job [years] 21	Satisfaction with employer Happy	Gender male	Age 53	Country New Guinea	Size of City 1316698	Profession probation officer	Unive De
795959	739095	1999.0	Large Apartment	33	8	Happy	unknown	17	Togo	92767	industrial program compliance analyst	
180758	180759	1954.0	nA	110	10	Somewhat Happy	female	21	Portugal	1296332	office administrator	M:
175084	175085	1954.0	nA	135	24	Average	male	58	Czechia	1426606	scientistscientist	Bac

In []:

In [10]:

```
data = pd.concat([train,test],ignore_index=True)

data['University Degree']=data['University Degree'].fillna('Bachelor')

data['Gender']=data['Gender'].replace('m','male')
data['Gender']=data['Gender'].replace('f','female')
data['Gender']=data['Gender'].replace('unknown','other')
data['Gender']=data['Gender'].fillna('female')

#data['Housing Situation']=data['Housing Situation'].replace('nA','0')
data['Housing Situation']=np.where(data['Housing Situation']=='0', 'nA', data['Housing Situation'])
data['Housing Situation']=np.where(data['Housing Situation']==0, 'nA', data['Housing Situation'])
data['Hair Color']=data['Hair Color'].fillna(method='bfill')

data['Satisfaction with employer']=data['Satisfaction with employer'].fillna('Average')

data.fillna(value={'Year of Record':data['Year of Record'].mean()}, inplace=True)

data['Profession']=data['Profession'].fillna(method='bfill')

data['Country']=data['Country'].fillna(method='bfill')

data.shape
```

Out[10]:

(1418012, 17)

In [11]:

```
#构造等级特征
data['Satisfaction with employer'] = data['Satisfaction with employer'].map \
    ({'Average':2, 'Happy':4, 'Somewhat Happy':3, 'Unhappy':1})
```

In [12]:

```
data.isnull().any()
```

Out[12]:

Instance	False
Year of Record	False
Housing Situation	False
Crime Level in the City of Employment	False
Work Experience in Current Job [years]	False
Satisfaction with employer	False
Gender	False
Age	False
Country	False
Size of City	False

```

Profession                False
University Degree          False
Wears Glasses             False
Hair Color                False
Body Height [cm]          False
Yearly Income in addition to Salary (e.g. Rental Income) False
Total Yearly Income [EUR] True
dtype: bool

```

In [13]:

```

#对于每个country和profession特征, 用其特征值下收入均值来替换
country_income = dict(train.groupby('Country').mean()['Total Yearly Income [EUR]']/10000)
data.Country = data.Country.map(country_income)
data.Country = data.Country.fillna(data.Country.mean())
country_income = dict(train.groupby('Profession').mean()['Total Yearly Income [EUR]']/10000)
data.Profession = data.Profession.map(country_income)
country_income = dict(train.groupby('Profession').mean()['Total Yearly Income [EUR]']/10000)
data.Profession = data.Profession.map(country_income)
#前面的254287数据用来构造均值特征
sp = 254287

```

In [14]:

```

#转换成数值
data.iloc[:, -2] = data.iloc[:, -2].map(lambda x: float(x[:-3]))

```

In [15]:

```

data['BigCity'] = np.where(data['Size of City']>7335190, 1, 0)
data['SmallCity'] = np.where(data['Size of City']<7335190, 1, 0)
#data = data.drop(columns=["Size of City"])

```

In [16]:

```

data['Crime Level in the City of Employment']=data['Crime Level in the City of
Employment'].replace(0,data['Crime Level in the City of Employment'].mean())

```

In [17]:

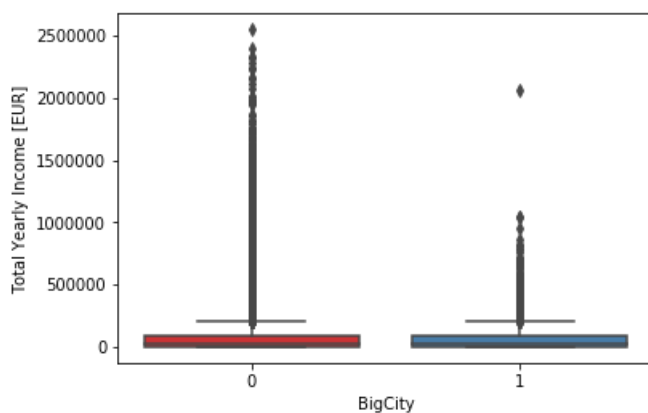
```

sns.boxplot(x=data['BigCity'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")

```

Out[17]:

<matplotlib.axes._subplots.AxesSubplot at 0x11e501c50>



In []:

```

# Remove outliers in Size of City
#indexBigCityOutliers = data[ (data["BigCity"] == 1) & (data["Total Yearly Income [EUR]"] > 1500000) ].index
#indexBigCityOutliers

```

In []:

```
#data = data.drop(indexBigCityOutliers)
```

In [18]:

```
data['Housing Situation'].value_counts()
```

Out[18]:

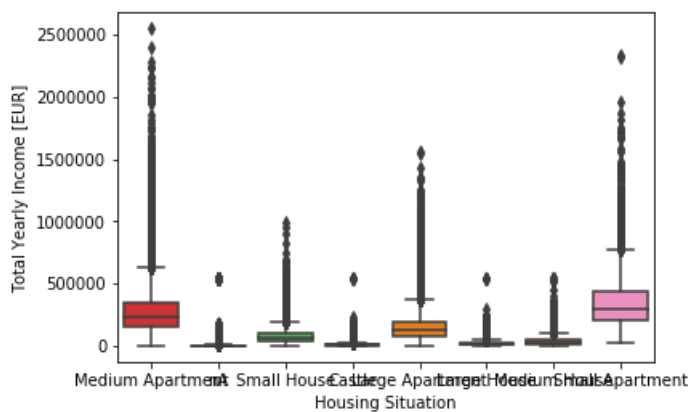
```
nA                365345
Large House       212200
Medium House      184710
Castle            170926
Large Apartment   170623
Small House       169785
Medium Apartment  134157
Small Apartment   10266
Name: Housing Situation, dtype: int64
```

In [19]:

```
# Have changed 0's and '0' to nA
sns.boxplot(x=data['Housing Situation'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x11d01d4a8>

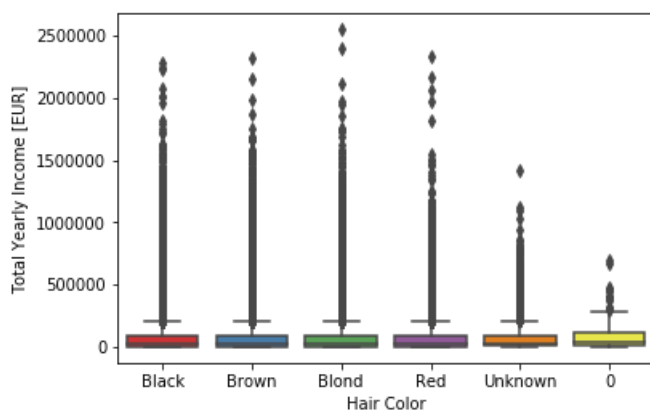


In [20]:

```
sns.boxplot(x=data['Hair Color'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[20]:

<matplotlib.axes._subplots.AxesSubplot at 0x11cc5cda0>

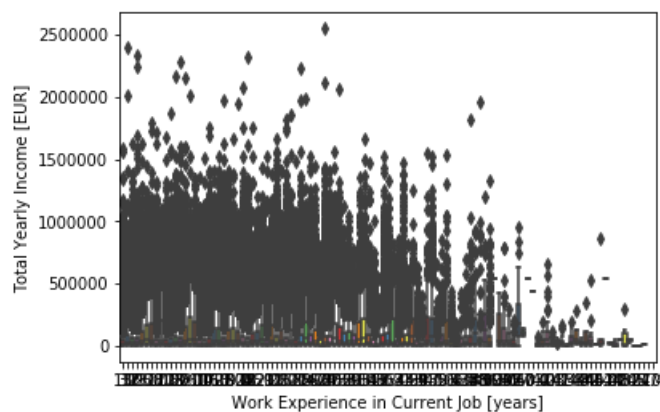


In [30]:

```
sns.boxplot(x=data['Work Experience in Current Job [years]'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[30]:

<matplotlib.axes._subplots.AxesSubplot at 0x123cc2828>




In [31]:

```
sns.boxplot(x=data['University Degree'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[31]:

<matplotlib.axes._subplots.AxesSubplot at 0x160e26b70>

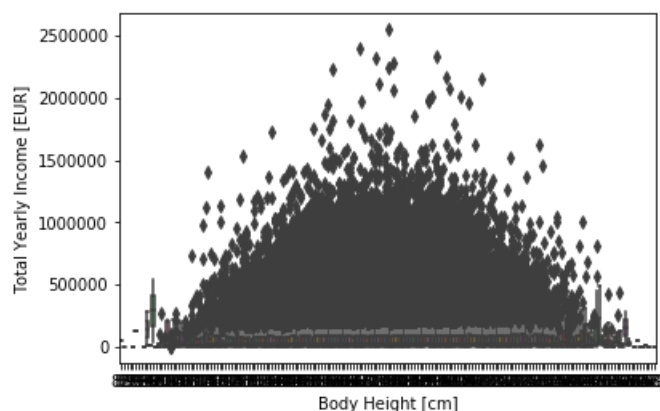


In [32]:

```
sns.boxplot(x=data['Body Height [cm]'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[32]:

<matplotlib.axes._subplots.AxesSubplot at 0x150c91860>

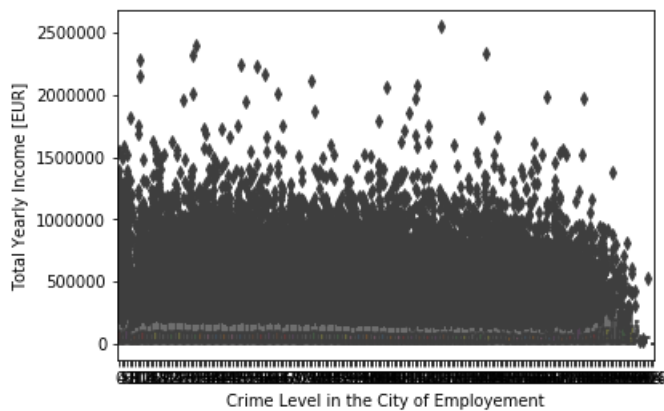


In [33]:

```
sns.boxplot(x=data['Crime Level in the City of Employment'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[33]:

<matplotlib.axes._subplots.AxesSubplot at 0x137e39c88>

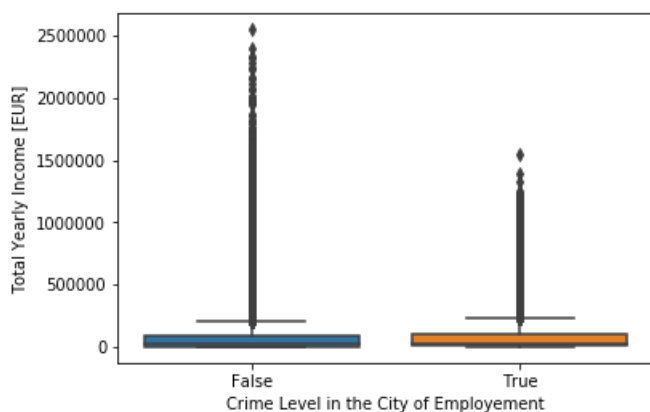


In [45]:

```
sns.boxplot(x = data['Crime Level in the City of Employment']==0,y=data["Total Yearly Income [EUR]"])
```

Out[45]:

<matplotlib.axes._subplots.AxesSubplot at 0x14da809b0>




In [46]:

```
sns.boxplot(x=data['Gender'], y=data["Total Yearly Income [EUR]"], data=data, palette="Set1")
```

Out[46]:

<matplotlib.axes._subplots.AxesSubplot at 0x14d5a1f60>



In []:

In []:

```
data.head()
```

In [21]:

```
cats = ['Year of Record', 'Housing Situation', 'Country', 'Size of City',
        'Crime Level in the City of Employment', 'Work Experience in Current Job [years]']
cons = ['Satisfaction with employer', 'Gender', 'Age',
        'University Degree', 'Body Height [cm]', 'Profession']
data['Work Experience in Current Job [years]'] = data['Work Experience in Current Job [years]'].as
type(str)
```

In [22]:

```
#This is the inspiration I got from the best code.
#I added and constructed mean features, cross mean features and Category Characteristics

def create_feature(df, cats, cons, normalize=True):
    for cat in cats:
        value = df[cat].value_counts(dropna=False, normalize=normalize).to_dict()
        num = cat + '_FE_FULL'
        df[num + num] = df[cat].map(value)
        #构造均值特征mean features
        df[num] = df[cat].map(dict(df.iloc[:sp].groupby(cat).mean()['Total Yearly Income [EUR]']/1
0000))
        df[num] = df[num].fillna(df[num].mean())
        df[num] = df[num].astype('float32')
        for con in cons:
            new_col = cat + '_' + con
            df[new_col] = df[cat].astype(str) + '_' + df[con].astype(str)
            temp_df = df[new_col]
            fq_encode = temp_df.value_counts(normalize=True).to_dict()
            #构造交叉均值特征cross mean features
            df[new_col] = df[new_col].map(dict(df.iloc[:sp+1].groupby( \
                new_col).mean()['Total Yearly Income [EUR]']/10000))
            df[new_col] = df[new_col].fillna(df[new_col].mean())
    return df

data = create_feature(data, cats, cons)
data['Work Experience in Current Job [years]'] = data['Work Experience in Current Job [years]' \
                                                    ].replace('#NUM!', data.iloc[:, -1].mean()).as
pe(float)

#构造类别特征 Category Characteristics
for col in data.dtypes[data.dtypes == 'object'].index.tolist():
    feat_le = LabelEncoder()
    feat_le.fit(data[col].unique().astype(str))
    data[col] = feat_le.transform(data[col].astype(str))

del_col = set(['Total Yearly Income [EUR]', 'Instance'])
features_col = list(set(data) - del_col)
features_col
```

Out[22]:

```
['Satisfaction with employer',
 'Profession',
 'Crime Level in the City of Employment_Age',
 'Year of Record_FE_FULL',
 'Year of Record_Profession',
 'Size of City_Gender',
 'University Degree',
 'Year of Record_Satisfaction with employer',
 'Year of Record_Age',
 'Work Experience in Current Job [years]_Profession',
 'Housing Situation_Profession',
 'Housing Situation',
```

```

'Age',
'Yearly Income in addition to Salary (e.g. Rental Income)',
'Size of City_FE_FULLSize of City_FE_FULL',
'Year of Record_Gender',
'Year of Record_Body Height [cm]',
'Housing Situation_FE_FULLHousing Situation_FE_FULL',
'Country_FE_FULLCountry_FE_FULL',
'Size of City_Profession',
'Work Experience in Current Job [years]_Gender',
'Crime Level in the City of Employment_FE_FULL',
'Year of Record',
'Housing Situation_Body Height [cm]',
'Work Experience in Current Job [years]_Satisfaction with employer',
'Body Height [cm]',
'Work Experience in Current Job [years]_University Degree',
'Crime Level in the City of Employment',
'Housing Situation_FE_FULL',
'Size of City_Age',
'Gender',
'Size of City',
'Crime Level in the City of Employment_Gender',
'Country_FE_FULL',
'Size of City_Satisfaction with employer',
'Country',
'Housing Situation_Gender',
'Size of City_FE_FULL',
'SmallCity',
'Crime Level in the City of Employment_Profession',
'BigCity',
'Crime Level in the City of Employment_FE_FULLCrime Level in the City of Employment_FE_FULL',
'Work Experience in Current Job [years]_FE_FULL',
'Work Experience in Current Job [years]',
'Country_University Degree',
'Housing Situation_Age',
'Country_Satisfaction with employer',
'Crime Level in the City of Employment_Satisfaction with employer',
'Year of Record_University Degree',
'Wears Glasses',
'Country_Age',
'Size of City_Body Height [cm]',
'Work Experience in Current Job [years]_Age',
'Crime Level in the City of Employment_Body Height [cm]',
'Country_Body Height [cm]',
'Year of Record_FE_FULLYear of Record_FE_FULL',
'Country_Gender',
'Hair Color',
'Work Experience in Current Job [years]_FE_FULLWork Experience in Current Job [years]_FE_FULL',
'Housing Situation_University Degree',
'Country_Profession',
'Size of City_University Degree',
'Crime Level in the City of Employment_University Degree',
'Housing Situation_Satisfaction with employer',
'Work Experience in Current Job [years]_Body Height [cm]'

```

In [24]:

```
data.shape
```

Out[24]:

```
(1418012, 67)
```

In [25]:

```

from sklearn.ensemble import RandomForestRegressor

param = {'num_iterations':20000,
         'max_depth': 21,
         'objective':'regression',
         "verbosity": -1,
         'metric': 'mae',
         'bagging_fraction': 0.8,
         'learning_rate': 0.01,}

X_train,X_test = data[features_col].iloc[:1048573],data[features_col].iloc[1048574:]
Y_train = data['Total Yearly Income [EUR]'].iloc[:1048573]

```



```

X_train = data[total_really_income [BANK]].iloc[:10000]
x_train,x_val,y_train,y_val = X_train.iloc[sp+1:,:], X_train.iloc[:sp,:], \
    Y_train.iloc[sp+1:], Y_train.iloc[:sp]
train_data = lgb.Dataset(x_train, label=y_train, feature_name='auto')#categorical_feature=cat
val_data = lgb.Dataset(x_val, label=y_val, feature_name='auto')

bst = lgb.train(param, train_data, 20000, verbose_eval = 100, valid_sets=[val_data])

```

```

/Users/BarryFitzpatrick/anaconda3/lib/python3.7/site-packages/lightgbm/engine.py:148: UserWarning:
Found `num_trees` in params. Will use it instead of argument
  warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))

```

```

[100] valid_0's l1: 31607.9
[200] valid_0's l1: 18388.6
[300] valid_0's l1: 13825.5
[400] valid_0's l1: 12009.1
[500] valid_0's l1: 11184.9
[600] valid_0's l1: 10777.6
[700] valid_0's l1: 10508.3
[800] valid_0's l1: 10305.7
[900] valid_0's l1: 10144.1
[1000] valid_0's l1: 9990.41
[1100] valid_0's l1: 9876.47
[1200] valid_0's l1: 9786.57
[1300] valid_0's l1: 9684.97
[1400] valid_0's l1: 9606.25
[1500] valid_0's l1: 9539.65
[1600] valid_0's l1: 9479.77
[1700] valid_0's l1: 9431.6
[1800] valid_0's l1: 9380.01
[1900] valid_0's l1: 9331.96
[2000] valid_0's l1: 9294.14
[2100] valid_0's l1: 9257.87
[2200] valid_0's l1: 9222.84
[2300] valid_0's l1: 9186.14
[2400] valid_0's l1: 9158.62
[2500] valid_0's l1: 9131.94
[2600] valid_0's l1: 9108.56
[2700] valid_0's l1: 9086.61
[2800] valid_0's l1: 9068
[2900] valid_0's l1: 9049.9
[3000] valid_0's l1: 9033.42
[3100] valid_0's l1: 9017.55
[3200] valid_0's l1: 9002.63
[3300] valid_0's l1: 8984.13
[3400] valid_0's l1: 8966.84
[3500] valid_0's l1: 8948.6
[3600] valid_0's l1: 8932.73
[3700] valid_0's l1: 8914.92
[3800] valid_0's l1: 8898.85
[3900] valid_0's l1: 8882.64
[4000] valid_0's l1: 8866.11
[4100] valid_0's l1: 8852.25
[4200] valid_0's l1: 8838.08
[4300] valid_0's l1: 8823.32
[4400] valid_0's l1: 8812.01
[4500] valid_0's l1: 8798.16
[4600] valid_0's l1: 8785.47
[4700] valid_0's l1: 8772.14
[4800] valid_0's l1: 8761.49
[4900] valid_0's l1: 8754.27
[5000] valid_0's l1: 8744.15
[5100] valid_0's l1: 8734.49
[5200] valid_0's l1: 8724.68
[5300] valid_0's l1: 8713.63
[5400] valid_0's l1: 8706.56
[5500] valid_0's l1: 8697.85
[5600] valid_0's l1: 8688.99
[5700] valid_0's l1: 8682.59
[5800] valid_0's l1: 8675.27
[5900] valid_0's l1: 8670.71
[6000] valid_0's l1: 8665.22
[6100] valid_0's l1: 8656.41
[6200] valid_0's l1: 8649.41
[6300] valid_0's l1: 8638.32
[6400] valid_0's l1: 8626.14

```

[6500] valid_0's l1: 8614.37
[6600] valid_0's l1: 8605.72
[6700] valid_0's l1: 8593.99
[6800] valid_0's l1: 8585.28
[6900] valid_0's l1: 8578.78
[7000] valid_0's l1: 8570.31
[7100] valid_0's l1: 8562.18
[7200] valid_0's l1: 8555.06
[7300] valid_0's l1: 8548.43
[7400] valid_0's l1: 8539.6
[7500] valid_0's l1: 8537.19
[7600] valid_0's l1: 8534.81
[7700] valid_0's l1: 8531.22
[7800] valid_0's l1: 8526.3
[7900] valid_0's l1: 8521.84
[8000] valid_0's l1: 8516.86
[8100] valid_0's l1: 8512.01
[8200] valid_0's l1: 8507.82
[8300] valid_0's l1: 8503.25
[8400] valid_0's l1: 8497.41
[8500] valid_0's l1: 8492.74
[8600] valid_0's l1: 8488.96
[8700] valid_0's l1: 8485.59
[8800] valid_0's l1: 8482.55
[8900] valid_0's l1: 8479.56
[9000] valid_0's l1: 8476.23
[9100] valid_0's l1: 8473.74
[9200] valid_0's l1: 8469.83
[9300] valid_0's l1: 8466.43
[9400] valid_0's l1: 8464.56
[9500] valid_0's l1: 8460.63
[9600] valid_0's l1: 8456.23
[9700] valid_0's l1: 8450.7
[9800] valid_0's l1: 8445.39
[9900] valid_0's l1: 8440.4
[10000] valid_0's l1: 8436.14
[10100] valid_0's l1: 8432.02
[10200] valid_0's l1: 8429.33
[10300] valid_0's l1: 8425.67
[10400] valid_0's l1: 8423.09
[10500] valid_0's l1: 8420.76
[10600] valid_0's l1: 8418.82
[10700] valid_0's l1: 8413.58
[10800] valid_0's l1: 8410.69
[10900] valid_0's l1: 8406.74
[11000] valid_0's l1: 8406
[11100] valid_0's l1: 8402.51
[11200] valid_0's l1: 8399.41
[11300] valid_0's l1: 8397.48
[11400] valid_0's l1: 8396
[11500] valid_0's l1: 8393.06
[11600] valid_0's l1: 8388.99
[11700] valid_0's l1: 8385
[11800] valid_0's l1: 8381.95
[11900] valid_0's l1: 8380.93
[12000] valid_0's l1: 8377.49
[12100] valid_0's l1: 8375.03
[12200] valid_0's l1: 8373.55
[12300] valid_0's l1: 8370.76
[12400] valid_0's l1: 8369.19
[12500] valid_0's l1: 8365.68
[12600] valid_0's l1: 8363.84
[12700] valid_0's l1: 8361.54
[12800] valid_0's l1: 8359.56
[12900] valid_0's l1: 8358.96
[13000] valid_0's l1: 8357.79
[13100] valid_0's l1: 8355.4
[13200] valid_0's l1: 8355.11
[13300] valid_0's l1: 8352.24
[13400] valid_0's l1: 8351.7
[13500] valid_0's l1: 8349.29
[13600] valid_0's l1: 8347.28
[13700] valid_0's l1: 8344.46
[13800] valid_0's l1: 8342.26
[13900] valid_0's l1: 8340.55
[14000] valid_0's l1: 8337.87
[14100] valid_0's l1: 8336.65

```
[14200] valid_0's l1: 8334.61
[14300] valid_0's l1: 8332.52
[14400] valid_0's l1: 8331.24
[14500] valid_0's l1: 8326.27
[14600] valid_0's l1: 8320
[14700] valid_0's l1: 8319.8
[14800] valid_0's l1: 8318.77
[14900] valid_0's l1: 8318.31
[15000] valid_0's l1: 8317.23
[15100] valid_0's l1: 8315.83
[15200] valid_0's l1: 8314.7
[15300] valid_0's l1: 8310.9
[15400] valid_0's l1: 8310.24
[15500] valid_0's l1: 8305.88
[15600] valid_0's l1: 8302.13
[15700] valid_0's l1: 8299.68
[15800] valid_0's l1: 8299.91
[15900] valid_0's l1: 8298.55
[16000] valid_0's l1: 8296.96
[16100] valid_0's l1: 8292.89
[16200] valid_0's l1: 8290.75
[16300] valid_0's l1: 8290.33
[16400] valid_0's l1: 8289.35
[16500] valid_0's l1: 8289.21
[16600] valid_0's l1: 8285.32
[16700] valid_0's l1: 8282.37
[16800] valid_0's l1: 8280.04
[16900] valid_0's l1: 8276.58
[17000] valid_0's l1: 8275.6
[17100] valid_0's l1: 8273.91
[17200] valid_0's l1: 8271.35
[17300] valid_0's l1: 8267.61
[17400] valid_0's l1: 8265.11
[17500] valid_0's l1: 8263.88
[17600] valid_0's l1: 8262.39
[17700] valid_0's l1: 8259.87
[17800] valid_0's l1: 8258.94
[17900] valid_0's l1: 8255.94
[18000] valid_0's l1: 8254.23
[18100] valid_0's l1: 8251.35
[18200] valid_0's l1: 8249.55
[18300] valid_0's l1: 8248.63
[18400] valid_0's l1: 8246.3
[18500] valid_0's l1: 8244.57
[18600] valid_0's l1: 8243.39
[18700] valid_0's l1: 8241.38
[18800] valid_0's l1: 8239
[18900] valid_0's l1: 8238.07
[19000] valid_0's l1: 8236.23
[19100] valid_0's l1: 8234.59
[19200] valid_0's l1: 8232.12
[19300] valid_0's l1: 8230.02
[19400] valid_0's l1: 8228.55
[19500] valid_0's l1: 8227.63
[19600] valid_0's l1: 8226.7
[19700] valid_0's l1: 8225.29
[19800] valid_0's l1: 8224.55
[19900] valid_0's l1: 8223.83
[20000] valid_0's l1: 8222.01
```

In []:

```
from sklearn.metrics import mean_absolute_error
predict = bst.predict(x_val)
val_mae = mean_absolute_error(y_val,predict)
val_mae
```

In []:

```
#生成结果
#rfr.fit(X_train, Y_train)
predict = bst.predict(X_test)
result=pd.DataFrame([range(1,1+len(predict)), predict]).T
result.columns = ['Instance', 'Total Yearly Income [EUR]']
result.to_csv("sub191125_10.csv",index=False)
```

```
result.head()
```

```
In [ ]:
```